

Suhaas HW 3 ISLR

Suhaas Adiraju

2024-10-01

Homework 3

Question 8

```
library(ISLR) ### a
```

```
data(USArrests)
states = row.names(USArrests)
names(USArrests)
```

```
## [1] "Murder" "Assault" "UrbanPop" "Rape"
```

```
dim(USArrests)
```

```
## [1] 50 4
```

```
apply(USArrests, 2, mean)
```

```
## Murder Assault UrbanPop Rape
## 7.788 170.760 65.540 21.232
```

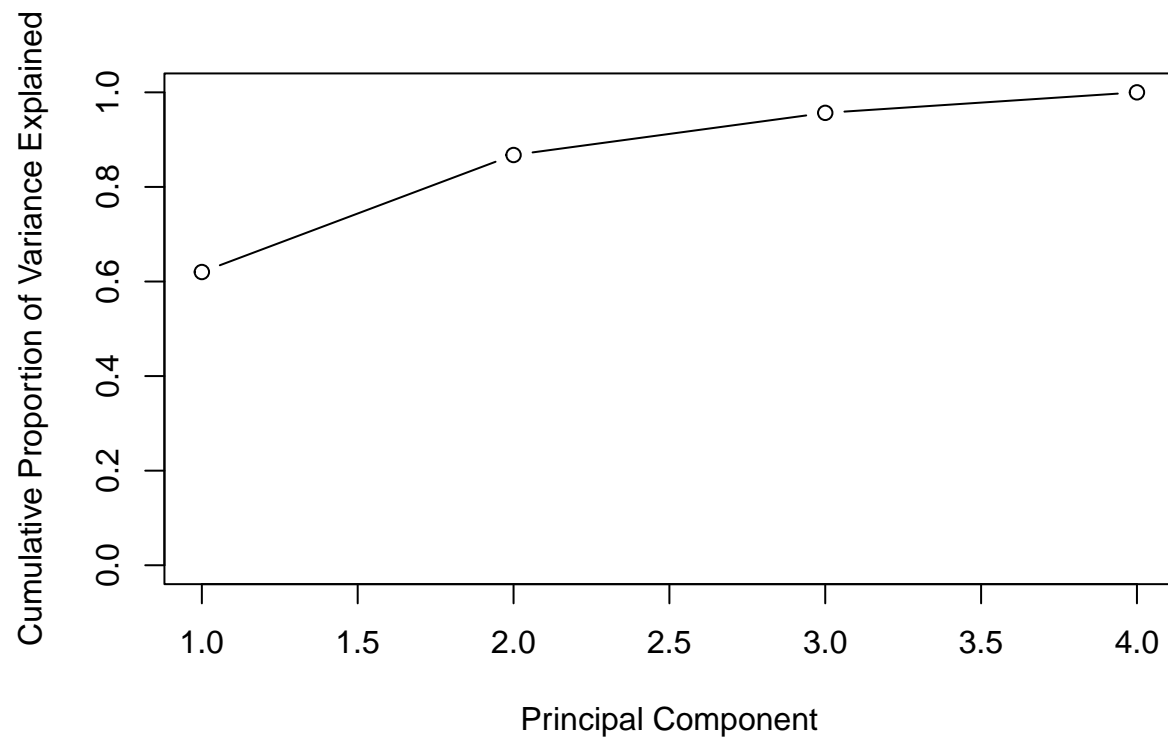
```
apply(USArrests, 2, var)
```

```
## Murder Assault UrbanPop Rape
## 18.97047 6945.16571 209.51878 87.72916
```

```
pr.out = prcomp(USArrests, scale=TRUE)
pr.var = pr.out$sdev^2
PVE = pr.var/sum(pr.var)
PVE
```

```
## [1] 0.62006039 0.24744129 0.08914080 0.04335752
```

```
plot(cumsum(PVE), xlab="Principal Component", ylab = "
Cumulative Proportion of Variance Explained", ylim=c(0,1),
type='b')
```



b

```
pr.loadings = pr.out$rotation
PC1var = var(pr.out$x[,1])
PC2var = var(pr.out$x[,2])
PC3var = var(pr.out$x[,3])
PC4var = var(pr.out$x[,4])

totalVar = (PC1var+PC2var+PC3var+PC4var)

PC1.PVEhand = PC1var/totalVar
PC2.PVEhand = PC2var/totalVar
PC3.PVEhand = PC3var/totalVar
PC4.PVEhand = PC4var/totalVar
PC1.PVEhand
```

```
## [1] 0.6200604
```

```
PC2.PVEhand
```

```
## [1] 0.2474413
```

```
PC3.PVEhand
```

```
## [1] 0.0891408
```

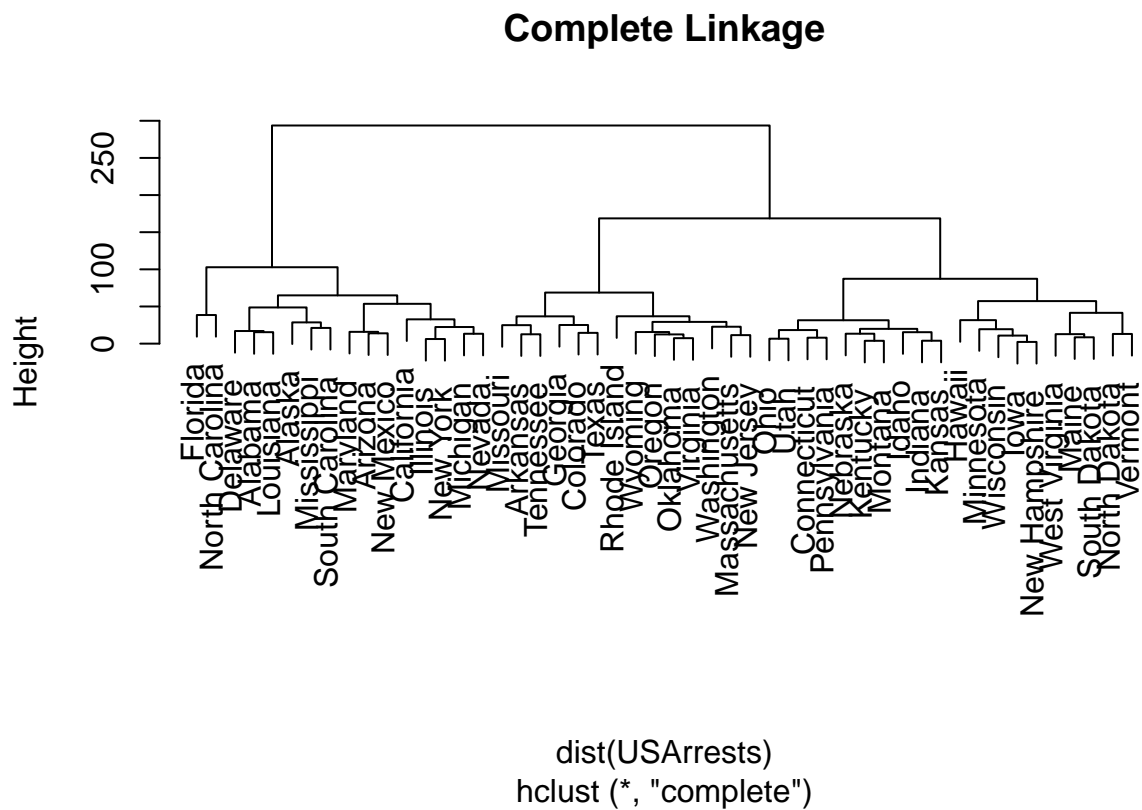
```
PC4.PVEhand
```

```
## [1] 0.04335752
```

Question 9

a

```
Arr.complete = hclust(dist(USArrests),method='complete')  
plot(Arr.complete,main="Complete Linkage")
```



b

```
print(sort(cutree(Arr.complete,3),decreasing=TRUE))
```

```
##      Connecticut      Hawaii      Idaho      Indiana      Iowa
##          3          3          3          3          3
##      Kansas      Kentucky      Maine      Minnesota      Montana
##          3          3          3          3          3
##      Nebraska New Hampshire North Dakota      Ohio      Pennsylvania
##          3          3          3          3          3
##      South Dakota      Utah      Vermont West Virginia      Wisconsin
##          3          3          3          3          3
##      Arkansas      Colorado      Georgia Massachusetts      Missouri
##          2          2          2          2          2
##      New Jersey      Oklahoma      Oregon      Rhode Island      Tennessee
##          2          2          2          2          2
##      Texas      Virginia      Washington      Wyoming      Alabama
##          2          2          2          2          1
##      Alaska      Arizona      California      Delaware      Florida
##          1          1          1          1          1
##      Illinois      Louisiana      Maryland      Michigan      Mississippi
##          1          1          1          1          1
##      Nevada      New Mexico      New York North Carolina South Carolina
##          1          1          1          1          1
```

c

```
Arrest.scaled = scale(USArrests)
sd(Arrest.scaled[,1])
```

```
## [1] 1
```

```
sd(Arrest.scaled[,2])
```

```
## [1] 1
```

```
sd(Arrest.scaled[,3])
```


```
## [1] 1
```

```
sd(Arrest.scaled[,4])
```

```
## [1] 1
```

```
Arr.complete.scaled= hclust(dist(Arrest.scaled),method='complete')
plot(Arr.complete.scaled,main="Complete Linkage")
```

Complete Linkage



```
dist(Arrest.scaled)
hclust (*, "complete")
```

```
print(sort(cutree(Arr.complete.scaled,3),decreasing=TRUE))
```

```
##      Arkansas      Connecticut      Delaware      Hawaii      Idaho
##           3              3              3              3              3
##      Indiana          Iowa          Kansas      Kentucky      Maine
##           3              3              3              3              3
## Massachusetts      Minnesota      Missouri      Montana      Nebraska
##           3              3              3              3              3
## New Hampshire      New Jersey      North Dakota      Ohio      Oklahoma
##           3              3              3              3              3
##      Oregon      Pennsylvania      Rhode Island      South Dakota      Utah
##           3              3              3              3              3
##      Vermont      Virginia      Washington      West Virginia      Wisconsin
##           3              3              3              3              3
##      Wyoming      Arizona      California      Colorado      Florida
##           3              2              2              2              2
##      Illinois      Maryland      Michigan      Nevada      New Mexico
##           2              2              2              2              2
##      New York      Texas      Alabama      Alaska      Georgia
##           2              2              1              1              1
##      Louisiana      Mississippi      North Carolina      South Carolina      Tennessee
##           1              1              1              1              1
```

d

```
threeclust = sum(cutree(Arr.complete,3)==3)
threeclustScale = sum(cutree(Arr.complete.scaled,3)==3)
twoclust = sum(cutree(Arr.complete,3)==2)
twoclustScale = sum(cutree(Arr.complete.scaled,3)==2)
oneclust = sum(cutree(Arr.complete,3)==1)
oneclustScale = sum(cutree(Arr.complete.scaled,3)==1)

print('3 groups raw vs scaled:')
```

```
## [1] "3 groups raw vs scaled:"
```

```
threeclust
```

```
## [1] 20
```

```
threeclustScale
```

```
## [1] 31
```

```
print('2 groups raw vs scaled:')
```

```
## [1] "2 groups raw vs scaled:"
```

```
twoclust
```

```
## [1] 14
```

```
twoclustScale
```

```
## [1] 11
```

```
print('1 group raw vs scaled:')
```

```
## [1] "1 group raw vs scaled:"
```

```
oneclust
```

```
## [1] 16
```

```
oneclustScale
```

```
## [1] 8
```

We can see using the cutree function, that prior to scaling, the clustering includes less states in larger clusters, and more states in smaller clusters. This likely is a function of the disproportionate variability in the unscaled dataset thus more states are more unique, and thereby more singleton clusters are created. Disproportionate variability in unscaled data can be driven by high values in certain categories, for example urban population in california will be extremely high, and given we are using euclidean distance will be far away from other states if unscaled, this has a disproportionate effect on the state's position wrt arrests, but is not what we are interested in. So yes, I would scale data, IN THIS CASE. in general, the problem is very contextual. For example, if we simply had number of arrests per state, with no other variables included, I perhaps would not need to scale.

We can see in unscaled data california is in a singleton cluster, after scaling it is in a 2 state cluster. If we look at the crime values for california and compare to Nevada, whome it was most similarly clustered with, the crime values are highly similar

```
USArrests['California',]
```

```
##           Murder Assault UrbanPop Rape
## California      9      276      91 40.6
```

```
USArrests['Nevada',]
```

```
##           Murder Assault UrbanPop Rape
## Nevada    12.2      252      81  46
```

```
USArrests['Arizona',]
```

```
##           Murder Assault UrbanPop Rape
## Arizona     8.1      294      80  31
```

versus in the unscaled dendrogram, height wise it is similarly singleton to missouri or georgia. Which across all values seem much less similar to values seen in California

```
USArrests['California',]
```

```
##           Murder Assault UrbanPop Rape
## California      9      276      91 40.6
```

```
USArrests['Georgia',]
```

```
##           Murder Assault UrbanPop Rape
## Georgia    17.4      211      60 25.8
```

```
USArrests['Missouri',]
```

```
##           Murder Assault UrbanPop Rape
## Missouri      9      178      70 28.2
```

Question 11

a


```
getwd()
```

```
## [1] "C:/Users/suhaas.adiraju/Desktop/Statistical ML/StatisticalMLCourse"
```

```
Ch.Data= read.csv('C:\\Users\\suhaas.adiraju\\Desktop\\Statistical ML\\StatisticalMLCourse\\Ch10Ex11.csv')
```


b

```
sample = names(Ch.Data)
cor.mat.genes= as.dist(1-cor((Ch.Data)))
plot(hclust(cor.mat.genes,method='complete'),main='Correlation-based complete linkage clustering')
```



```
cor.mat.genes= as.dist(1-cor((Ch.Data)))
plot(hclust(cor.mat.genes,method='single'),main='Correlation-based single linkage clustering')
```



Correlation-based single linkage clustering



cor.mat.genes
hclust (*, "single")

```
cor.mat.genes= as.dist(1-cor((Ch.Data)))
plot(hclust(cor.mat.genes,method='average'),main='Correlation-based average linkage clustering')
```

Correlation-based average linkage clustering



```
cor.mat.genes
hclust(*, "average")
```

The separation of genes in to simply two groups is not trivial/automatic, but depends on the clustering used. none of the dendrograms only form two clusters, but using average and complete distances certainly provide a better 2 cluster estimation compared to single. Another way to show this result

```
sort(cutree(hclust(cor.mat.genes,method='average'),2))
```

```
## V1 V2 V3 V4 V7 V9 V16 V19 V20 V5 V6 V8 V10 V11 V12 V13 V14 V15 V17 V18
## 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2
## V21 V22 V23 V24 V25 V26 V27 V28 V29 V30 V31 V32 V33 V34 V35 V36 V37 V38 V39 V40
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
sort(cutree(hclust(cor.mat.genes,method='complete'),2))
```

```
## V1 V3 V4 V6 V8 V11 V12 V14 V15 V20 V2 V5 V7 V9 V10 V13 V16 V17 V18 V19
## 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2
## V21 V22 V23 V24 V25 V26 V27 V28 V29 V30 V31 V32 V33 V34 V35 V36 V37 V38 V39 V40
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
sort(cutree(hclust(cor.mat.genes,method='single'),2))
```

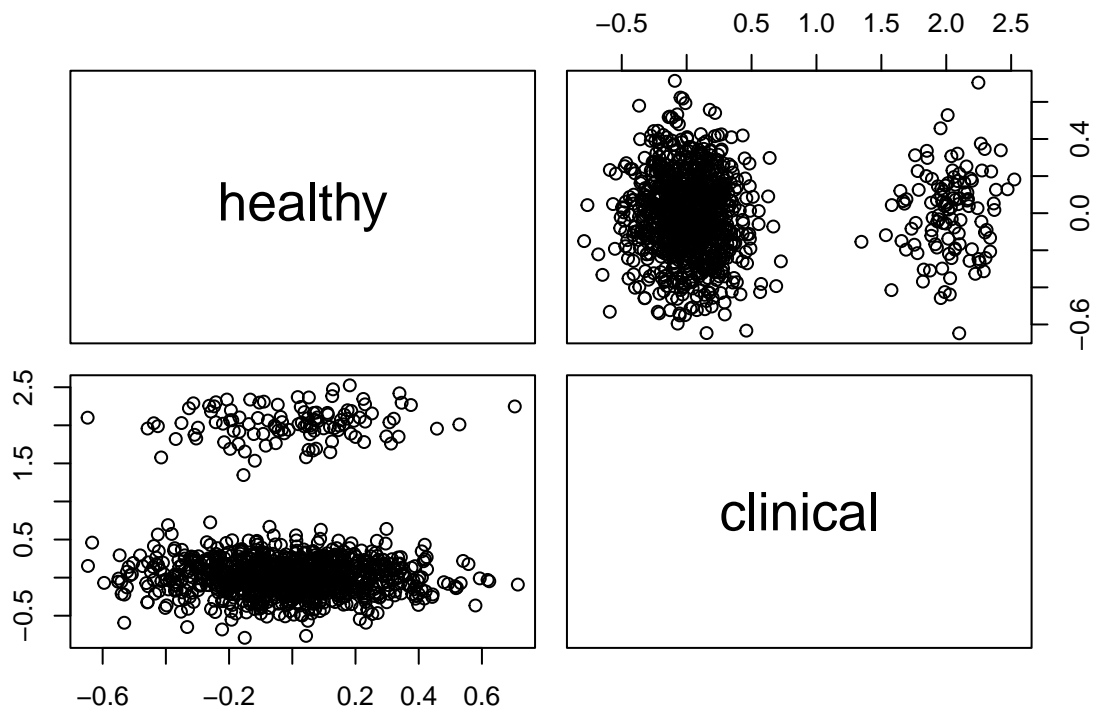
```
## V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14 V15 V16 V17 V18 V20 V21
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## V22 V23 V24 V25 V26 V27 V28 V29 V30 V31 V32 V33 V34 V35 V36 V37 V38 V39 V40 V19
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2
```

We can see single linkage clustered every sample except for V19 in cluster 1, this is clearly not useful.

c

There are several ways we can explore this question, I would start with very basic exploratory analysis and plotting. We know our two groups already, this is an advantage, we can average each gene(row) across samples, within each respective group of healthy versus diseased. Then use the pairs function to see if there are major differences plotting healthy versus diseased gene average values:

```
healthy = apply(Ch.Data[,1:20],1,mean)
clinical = apply(Ch.Data[,21:40],1,mean)
gene.frame = data.frame(healthy,clinical)
pairs(gene.frame)
```



Look at that! it seems that indeed there is a subgroup of genes that express much higher mean values in the clinical group. We could pull out the identity of these genes by setting a threshold and indexing the data frame

```
clinical.hit.genes= which((gene.frame$clinical>1.0),useNames = TRUE)
clinical.hit.genes
```


```
## [1] 11 12 13 14 15 16 17 18 19 20 501 502 503 504 505 506 507 508
## [19] 509 510 511 512 513 514 515 516 517 518 519 520 521 522 523 524 525 526
## [37] 527 528 529 530 531 532 533 534 535 536 537 538 539 540 541 542 543 544
## [55] 545 546 547 548 549 550 551 552 553 554 555 556 557 558 559 560 561 562
## [73] 563 564 565 566 567 568 569 570 571 572 573 574 575 576 577 578 579 580
## [91] 581 582 583 584 585 586 587 588 589 590 591 592 593 594 595 596 597 598
## [109] 599 600
```

although this is useful information, when we average across samples in each group, we are assuming a lot of stability across samples or rather overlooking the within group variability, as well as other nuances. To take a more fine-grained approach of preserving individual sample contributions lets use PCA

```
pr.out.gene = prcomp((Ch.Data), scale=TRUE)
pr.var.gene = pr.out.gene$sdev^2
PVE.gene = pr.var.gene/sum(pr.var.gene)
PVE.gene
```

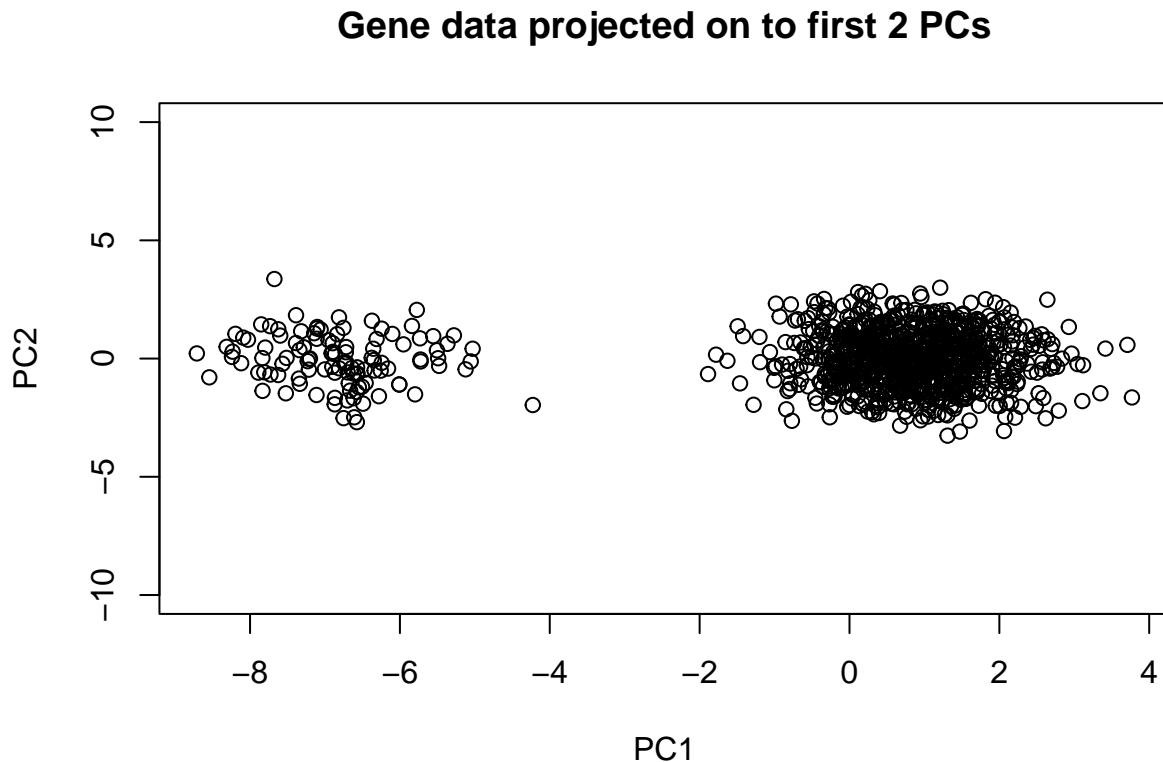
```
## [1] 0.16130137 0.03280179 0.03176830 0.03075558 0.02972109 0.02847757
## [7] 0.02833663 0.02783324 0.02736489 0.02674424 0.02610345 0.02519941
## [13] 0.02499759 0.02484029 0.02396715 0.02339585 0.02313948 0.02299186
## [19] 0.02250569 0.02193219 0.02091515 0.02041673 0.02023852 0.01921437
## [25] 0.01863679 0.01836530 0.01819977 0.01786939 0.01757227 0.01708818
## [31] 0.01658588 0.01587270 0.01569743 0.01524706 0.01516708 0.01485582
## [37] 0.01409545 0.01388385 0.01331786 0.01258275
```

```
plot(cumsum (PVE.gene ), xlab="Principal Component", ylab = "
Cumulative Proportion of Variance Explained", ylim=c(0,1) ,
)
```



We see from our cumulative plot and the proportion of variance explained values that our first component captures the highest proportion of variance, and all following components add about the same small amount of information. lets plot our data mapped in the first 2 components.

```
plot(pr.out.gene$x[, 'PC1'], pr.out.gene$x[, 'PC2'], ylim = c(-10, 10), main = 'Gene data projected on to first 2 PCs')
```



We can see in this plot the data varies a lot on PC1, and a little on PC2, and via the PC1 variance, there is a distinct data cloud with large negative score values, these could be positive too, as PCA weights are identical up to sign changes. Basically, these genes differ a lot among the samples. We should index them, and sort them, to see which genes they are and which are the most variable along PC1

```
gene.hits.pca.idx = which(pr.out.gene$x[, 'PC1'] <= -4, useNames = TRUE)
top.genes = data.frame(gene.hits.pca.idx, -(pr.out.gene$x[gene.hits.pca.idx, 'PC1']))
sorted.top.genes = order(top.genes$X.pr.out.gene.x.gene.hits.pca.idx...PC1..., decreasing = TRUE)
top.genes[sorted.top.genes, 1]
```

```
## [1] 551 589 508 548 509 511 566 565 540 534 561 502 584 586 12 558 593 568
## [19] 582 578 503 13 549 572 538 505 554 587 545 550 521 11 522 526 560 555
## [37] 570 597 559 541 516 590 571 576 517 14 546 600 520 529 515 514 506 535
## [55] 537 596 542 523 595 594 512 599 562 591 536 556 544 527 575 20 533 564
## [73] 530 539 579 18 513 543 567 524 501 569 583 563 588 531 16 585 592 581
## [91] 19 15 574 553 580 598 17 519 507 547 552 504 557 510 577 573 528 525
## [109] 532 518
```